

Computer Communications and Networks (COMN), 2016/17, Semester 2

Assignment

Overview

The overall goal of this assignment is to implement and evaluate different protocols for achieving end-to-end reliable data transfer at the application layer over the unreliable datagram protocol (UDP) transport protocol. In particular, you are asked to implement in Java three different sliding window protocols – *StopandWait*, *Go Back N* and *Selective Repeat* – at the application layer using UDP sockets. Note that the stopandwait protocol can be viewed as a special kind of sliding window protocol in which sender and receiver window sizes are both equal to 1. For each of the three sliding window protocols, you will implement the two protocol endpoints referred henceforth as *sender* and *receiver* respectively; these endpoints also act as application programs. Data communication is *unidirectional*, requiring transfer of a large file from the sender to the receiver over a link as a sequence of smaller messages. The underlying link is assumed to be *symmetric* in terms of bandwidth and delay characteristics.

To test your protocol implementations and study their performance in a controlled environment on DICE machines, you will need to use the *Dummynet* link emulator [1]. Specifically, the sender and receiver processes for each of the three protocols will run *within* the same (virtual) machine and communicate with each other over a link *emulated* by *Dummynet*. For this assignment, you only need the basic functionality of *Dummynet* to emulate a link with desired characteristics in terms of bandwidth, delay and packet loss rate.

Virtual Machine Setup

More specifically, you need to setup a virtual machine (VM) using your DICE accounts, following the instructions in [2], to run your protocol implementations and evaluate their performance. The VM so created (*DummynetSL6*) can be used from any DICE machine and comes with *Dummynet* preinstalled. You will be able to configure *Dummynet* using `ipfw` commands. More on this shortly.

Since *DummynetSL6* VM does not include *Eclipse*, we suggest you develop your protocol implementations outside it and save them within the *dummynetshared* subdirectory of your assignment directory. That way, the files will be accessible from within the VM via mount command described under “Shared folder” in [2]. You should however be able to compile and run your code from inside the VM as the Java compiler (`javac`) and application launcher (`java`) are installed as part of the *dummynetSL6* VM.

You can use the `/work` space within the *dummynetSL6* VM for storing any temporary files you would like to keep across various executions of the VM.

Link Emulation using Dummynet

Once the above onetime VM setup part is done, you can configure and use the *Dummynet* to realize an emulated link between two communicating processes (e.g., your sender and receiver programs) inside the *DummynetSL6* VM. For example, to create a symmetric 1Mbps emulated link with 5ms oneway propagation delay (thus, 10ms in total considering both directions) and 0.5% packet loss rate for each direction (thus, 1% packet loss rate in total), you create two dummynet pipes for each direction and configure them as follows (as *root*):

```
% ipfw add pipe 100 in
% ipfw add pipe 200 out
% ipfw pipe 100 config delay 5ms plr 0.005 bw 1Mbits/s
% ipfw pipe 200 config delay 5ms plr 0.005 bw 1Mbits/s
```

You can verify this configuration by using the following commands:

```
% ipfw list
% ipfw pipe 100 show
% ipfw pipe 200 show
```

You can use the following command to flush all previous configuration rules:

```
% ipfw flush
```

Note that in the above, the pipe identifiers (100 and 200) are arbitrarily chosen. You could instead use different numbers and still get the same effect. If a configuration for a pipe needs to be updated, then you reissue the corresponding

“config” command with the modified value(s). For example, if you want the bandwidth for pipe 200 (corresponding to the outgoing direction of traffic) to be changed to 10Mbps instead, then you run the following command:

```
% ipfw pipe 200 config delay 5ms plr 0.005 bw 10Mbps/s
```

A few additional notes about *Dummynet* and *DummynetSL6* VM follow. **Note that we assume packets are not corrupted in transit (i.e., no bit errors) via the Dummynet emulated links, so there is no need to implement error detection functionality such as checksum at the endpoints.** Whole packets, however, can be lost over the emulated link as determined by the packet loss rate (plr) setting when configuring the emulated link using *Dummynet*. For more information on *Dummynet*, please refer to [1] and the [Dummynet website](#).

A final note on *Dummynet* is that a total round-trip propagation delay would be twice as large as the specified delay when both sender and receiver are in the same host that enforces the delay. More precisely, packets between the sender and receiver traverse each pipe twice because *Dummynet* sits between IP layer and a NIC device driver, and the packets go to the receiver after they first reach the NIC device driver.

```
% ipfw pipe 100 config delay 5ms plr 0.005 bw 1Mbps/s
% ipfw pipe 200 config delay 5ms plr 0.005 bw 1Mbps/s
```

Given that both sender and receiver are in the *DummynetSL6* VM, a total round-trip propagation delay is 20ms, not 10ms, in the above example. Each part of the assignment specification below states the *Dummynet* configuration parameters. As long as they are configured as described, you should be able to do the assignment.

Besides *Dummynet*, *DummynetSL6* VM has other networking utilities that you may find useful while working on this assignment. These include:

- [iperf](#)
- [thrulay](#)
- [netcat](#)
- [Wireshark](#)
- [tcpdump](#)

Note that these tools are explicitly mentioned so that you know they are available to use. Except for *iperf*, you are not required to use the rest of them for this assignment.

Detailed Assignment Specification

The assignment needs to be done in two parts. The second part builds on the first part. Each part is further divided into two subparts as detailed below.

Assignment Part 1

Part 1a: Basic framework (large file transmission under ideal conditions)

Implement sender and receiver endpoints for transferring a large file given at [3] from the sender to the receiver on localhost over UDP as a sequence of small messages with 1KB maximum payload (NB. 1KB = 1024 bytes) via *Dummynet* emulated link configured with **10Mbps bandwidth, 5ms oneway propagation delay and 0% packet loss rate** (i.e., no packet loss). **In the sender code, insert, at a minimum, a 10ms gap (i.e., sleep for 10ms) after each packet transmission.** The *Dummynet* sets a queue size of 50 as default (100 at a maximum). Because the sending rate from the sender is typically larger than the link speed (10Mbps) specified here, the queue is likely to overflow and hence packets losses are unavoidable. To allow the test of an ideal, reliable channel case, the 10ms gap is suggested. If packet losses continue to occur, increase the time gap. **Note that inserting the time gap is only for Part 1a. From Part 1b onwards, you should remove the sleeping part (i.e., without the time gap) from the sender.**

The exchanged data messages must also have a header part for the sender to include a 16bit message sequence number (for duplicate filtering at the receiver in the next and subsequent parts) and a bit/byte flag to indicate the last message (i.e., endoffile). Name the sender and receiver developed in this part as **Sender1a.java** and **Receiver1a.java** respectively.

Part 1b: Stop-and-Wait

Extend sender and receiver applications from *Part 1a* to implement a stop-and-wait protocol described in section 3.4.1 in [4], specifically rdt3.0. **[Hint:** You need two finite state machines (FSMs); one for rdt3.0 sender and the other for rdt3.0 receiver. While the sender FSM is presented in [4], there is no rdt3.0 receiver

FSM. The rdt3.0 receiver FSM is the rdt2.2 receiver FSM in [4]. Convince yourself why before you begin to implement the rdt3.0 protocol.] Call the resulting two applications Sender1b.java and Receiver1b.java respectively. This part requires you to define an acknowledgement message that the receiver will use to inform the sender about the receipt of a data message. Discarding duplicates at the receiver end using sequence numbers put in by the sender is also required. You can test the working of duplicate detection functionality in your implementation by using a small retransmission timeout on the sender side.

Using a **5% packet loss rate for each direction (i.e., pipe)** and rest of *Dummynet* emulated link configuration parameters as before (i.e., **10Mbps bandwidth and 5ms oneway propagation delay**), experiment with different retransmission timeouts and the corresponding number of retransmissions and throughput.

Tabulate your observations in the space provided under Question 1 in the results sheet for Part 1 provided at [5]. For this, your sender implementation should count the number of retransmissions and measure average throughput (in KB/s), which is defined as the ratio of file size (in KB) to the transfer time (in seconds). Transfer time in turn can be measured at the sender as the interval between first message transmission time and acknowledgement receipt time for last message. Before the sender application finishes and quits, print the average throughput value to the standard output.

Under Question 2 in [5], discuss the impact of retransmission timeout on number of retransmissions and throughput. Also indicate the optimal timeout value from communication efficiency viewpoint (i.e., the timeout that minimises the number of retransmissions).

Assignment Part 2

Part 2a: Go-Back-N

Extend Sender1b.java and Receiver1b.java from Part 1 to implement the GoBack-N protocol as described in section 3.4.3 of [4], by allowing the sender window size to be greater than 1. Name the sender and receiver implementations from this part as **Sender2a.java** and **Receiver2b.java** respectively.

Experiment with different window sizes at the sender (increasing in powers of 2 starting from 1) and **different oneway propagation delay values (5ms, 25ms and 100ms)** in the emulator. For the 5ms case, use the “optimal” value for the retransmission timeout identified from the previous part. Adjust the timeout

suitably for the other two cases. Across all these experiments, use the following values for the other emulated link parameters: **for each direction (i.e., pipe), 10Mbps bandwidth and 0.5% packet loss rate**. Tabulate your results under Question 1 and answer Question 2 in the results sheet for Part 2 provided at [6].

Part 2b: Selective Repeat

Extend `Sender2a.java` and `Receiver2a.java` to implement the selective repeat protocol as described in section 3.4.4 of [4]. Call the resulting two applications as **Sender2b.java** and **Receiver2b.java** respectively.

By configuring the *Dummynet* link with, **for each direction (i.e., pipe), 10Mbps bandwidth, 25ms oneway propagation delay and 0.5% packet loss rate**, experiment with different window size values and complete the table under Question 3 and answer Question 4 in [6].

As a part of this step, also carry out an equivalent experiment using *iperf* with TCP within the *dummynetSL6* VM, i.e., both *iperf* client and server running inside it. Use `-M` option in *iperf* to set the maximum segment size to 1KB and vary the TCP window sizes using the `-w` option. Note that *iperf* actually allocates twice the specified value, and uses the additional buffer for administrative purposes and internal kernel structures. But this is normal because effectively TCP uses the value specified as the window size for the session, which is the parameter to be varied in this experiment. You also need to specify the file to be transferred (i.e., the one given at [3]) as one (`-F` option) of the parameters to *iperf* on the client side. In addition, you should use `-t` option as well (refer to FAQ below for more details). Use the results of this experiment to complete the table under Question 5 and answer Question 6 in [6].

Implementation Guidelines

Your programs must adhere to the following standard with both sender and receiver application programs to be run inside the *DummynetSL6* VM:

- Sender programs must be named as specified below and must accept the following options from the command line:
`java SenderX localhost <Port> <Filename> [RetryTimeout] [WindowSize]`
where `<Port>` is the port number used by the corresponding receiver.
`<Filename>` is the file to transfer.

The `RetryTimeout` should be specified for parts 1b, 2a and 2b, whereas `WindowSize` option is only relevant for parts 2a and 2b.

For example: `java Sender1a localhost 54321 sfile`

- Receiver programs must be named as specified below and must accept the following options from the command line:

`java ReceiverX <Port> <Filename> [WindowSize]`

where `<Port>` is the port number which the receiver will use for receiving messages from the sender.

`<Filename>` is the name to use for the received file to save on local disk.

The `WindowSize` is relevant only for part 2b as it is implicitly equal to 1 for parts 1b and 2a.

For example: `java Receiver1a 54321 rfile`

- You can choose to have common files with functions used in different parts but you are required to submit such common files along with necessary documentation.
- You need to take appropriate measures for terminating your sender applications by considering cases where receiver finishes while sender keeps waiting for acknowledgements.
- Please use comments in your code!
- Please start each source file with the following comment line:
`/* Forename Surname MatriculationNumber */`
For example: `/* John Doe 1234567 */`

Submission

Submission deadlines for this assignment are as follows:

- **Part 1 due by 4pm on Friday, 17th February 2017:**

For Part 1, you must submit an electronic version of your implementations for Parts 1a and 1b (`Sender1a.java`, `Receiver1a.java`, `Sender1b.java`, `Receiver1b.java` and any common files), corresponding executables and completed results sheet [\[5\]](#) (as PDF). Use the following submit command:

```
submit comn 1 <directoryname>
```

- **Part 2 due by 4pm on Friday, 24th March 2017:**

For Part 2, you must submit an electronic version of your implementations for Parts 2a and 2b (`Sender2a.java`, `Receiver2a.java`, `Sender2b.java`,

Receiver2b.java and any common files), corresponding executables and completed results sheet [6] (as PDF). Use the following submit command:

```
submit comn 2 <directoryname>
```

No late submissions are allowed, except under extenuating circumstances as per the [School of Informatics policy on late submission of coursework](#).

You are expected to work on this assignment on your own. Or else, you will be committing plagiarism (see [School of Informatics guidelines on academic misconduct](#)).

Assessment

This assignment accounts for the whole of your coursework mark (or, 40% of the overall course mark). Distribution of marks among the different parts (as percentage of the coursework mark) is given below:

- Part 1 (30%)
 - Part 1a (10%)
 - Part 1b (20%)
- Part 2 (70%)
 - Part 2a (30%)
 - Part 2b (40%)

References

1. M. Carbone and L. Rizzo, "[Dummynet Revisited](#)," SIGCOMM Computer Communication Review, Vol. 40, No. 2, pp. 1220 Apr 2010.
2. [VirtualBox VM Setup Instructions](#)
3. [Test file](#)
4. J. F. Kurose and K. W. Ross, "Computer Networking: A TopDown Approach" (6th edition), Pearson Education, 2013.
5. [Part 1 Results Sheet](#)
6. [Part 2 Results Sheet](#)

FAQs

1. There is no /disk/scratch folder in student.login.inf.ed.ac.uk server.
A: The server is not a DICE machine. Please make sure that you log into a DICE machine.
2. When I try to compile and run my code within the virtual machine from the mnt/shared directory, the terminal responds with a "Permission Denied" message. I cannot run and test my code.
A: In that directory, you need root privilege. Run "su" and enter the root password.
3. "Operation not permitted" error when using ipfw command.
A: To run ipfw command, you need root privilege. Run "su" and enter the root password.
4. What is the root password for the dummynet VM?
A: It can be found in [VirtualBox VM Setup Instructions](#).
5. Are we allowed to use any libraries in our programs that would make the work easier, or do we have to code the sender and receivers from scratch?
A: Standard JAVA libraries are allowed. If your code doesn't get compiled due to the use of any non-standard libraries, your submission will not be marked.
6. What happens in case, Sender sends last packet, Receiver gets it, but the ACK gets lost? I assume that Receiver shuts down after last message has been obtained, but that means that in case of ACK being lost, Sender never terminates, because it tries to send the packet all the time waiting for the ACK, but that case will never happen, because Receiver is already closed.
A: The loss of the last ACK becomes an issue in the assignment because there is no connection establishment/termination mechanism like one in TCP. Such a mechanism will prevent this ambiguity. Note that its implementation is out of scope of this assignment.
Without the mechanism in your program, you can still terminate the receiver if the bit/byte flag is set, which makes the sender keep sending the unACKed packet. To address the issue, implement a mechanism that the sender gives up the transmission after say, 10 trials. But, make sure that retransmission

time and count (i.e., 10) for the last packet is not considered for counting total number of retransmissions.

7. When I ran iperf with -M option to set an MSS value, I received a warning message like "WARNING: attempt to set TCP maximum segment size to 1024, but got 536". Will this warning affect the transfer process?

A: This warning can be ignored. Tcpdump allows you to capture packet size and to confirm that the MSS is set correctly as configured.

8. When the following command is executed, the bytes transferred are always smaller than the file size.

```
iperf -c localhost -M 1KB -F test.jpg -w 16KB
```

Shouldn't the transferred data amount be the same as the size of test.jpg?

A: Two options (-F and -t) should be set together. With the two options, iperf is terminated on the expiration of the time (set by -t) or the completion of file transfer (set by -F) whichever comes first. Without setting -t option, the time duration is by default 10 seconds. Thus, if the file transfer is not finished in 10 seconds, the amount of transferred data is less than the actual file size. Hence, the time duration value should be large enough (e.g., 60 seconds). Note that when iperf calculates the transferred size, it takes into account both header and payload sizes. Thus, the transferred size reported by iperf is larger than the actual file size, which is absolutely normal.